


```
EXT_PORT = 2
```

```
if received_on_port != EXT_PORT: # Heartbeat
```

```
    return ([], [])
```

```

1 from state import flow_map, flow_id_to_backend_id, backends, backend_ip_map, cht
2 EXP_TIME = 10 * 1000
3 BACKEND_EXP_TIME = 3600000000 * 1000
4 EXT_PORT = 2
5
6 if a_packet_received:
7     flow_map.expire_all(now - EXP_TIME)
8     backend_ip_map.expire_all(now - BACKEND_EXP_TIME)
9
10 h3 = pop_header(tcpudp, on_mismatch=([], []))
11 h2 = pop_header(ipv4, on_mismatch=([], []))
12 h1 = pop_header(ether, on_mismatch=([], []))
13
14 assert a_packet_received
15 assert h1.type == 8 # 0x0800 == IPv4 in big endian
16 assert h2.npid == 6 or h2.npid == 17 # 6/17 -> TCP/UDP
17
18 if received_on_port == EXT_PORT: # Packet from the external network - client
19     packet_flow = LoadBalancedFlowc(h2.saddr, h2.daddr, h3.src_port, h3.dst_port, h2.npid)
20     alloc_flow_and_process_packet = False;
21     if flow_map.has(packet_flow):
22         flow_id = flow_map.get(packet_flow)
23         backend_id = flow_id_to_backend_id.get(flow_id)
24         if backend_ip_map.has_idx(backend_id):
25             flow_map.refresh_idx(flow_map.get(packet_flow), now)
26             backend = backends.get(backend_id)
27             return ([backend.nic],
28                     [ether(h1, saddr=..., daddr=backend.mac),
29                      ipv4(h2, cksum=..., daddr=backend.ip),
30                      tcpudp(h3)])
31         else:
32             flow_map.erase(packet_flow)
33             alloc_flow_and_process_packet = True

```

```

34 else:
35     alloc_flow_and_process_packet = True
36 if alloc_flow_and_process_packet:
37     if backend_ip_map.exists_with_cht(cht, _LoadBalancedFlow_hash(packet_flow)):
38         bknd = backend_ip_map.choose_with_cht(cht, _LoadBalancedFlow_hash(packet_flow))
39         if not flow_map.full():
40             idx = the_index_allocated
41             flow_map.add(packet_flow, idx, now)
42             flow_id_to_backend_id.set(idx, bknd)
43             backend = backends.get(bknd)
44             return ([backend.nic],
45                     [ether(h1, saddr=..., daddr=backend.mac),
46                      ipv4(h2, cksum=..., daddr=backend.ip),
47                      tcpudp(h3)])
48         else:
49             return ([], [])
50 else: # A heartbeat from a backend
51     bknd_addr = ip_addr(h2.saddr)
52     if backend_ip_map.has(bknd_addr):
53         backend_ip_map.refresh_idx(backend_ip_map.get(bknd_addr), now)
54     else:
55         if not backend_ip_map.full():
56             idx = the_index_allocated
57             backend_ip_map.add(bknd_addr, idx, now)
58             backends.set(idx, LoadBalancedBackendc(received_on_port, h1.saddr, h2.saddr))
59     return ([], [])

```





Full-stack

Pay-as-you-go




Push-button





Support for partial specifications





Push-button

Pay-as-you-go

Full-stack

Support for partial specifications

EXT_PORT = 2

```
if received_on_port != EXT_PORT: # Heartbeat
    return ([], [])
```

```
1 from state import flow_map, flow_id_to_backend_id, backends, backend_ip_map, cht
2 EXP_TIME = 10 * 1000
3 BACKEND_EXP_TIME = 3600000000 * 1000
4 EXT_PORT = 2
5
6 if a_packet_received:
7     flow_map.expire_all(now - EXP_TIME)
8     backend_ip_map.expire_all(now - BACKEND_EXP_TIME)
9
10 h3 = pop_header(tcpudp, on_mismatch=([], []))
11 h2 = pop_header(ipv4, on_mismatch=([], []))
12 h1 = pop_header(ether, on_mismatch=([], []))
13
14 assert a_packet_received
15 assert h1.type == 8 # 0x0800 == IPv4 in big endian
16 assert h2.npid == 6 or h2.npid == 17 # 6/17 -> TCP/UDP
17
18 if received_on_port == EXT_PORT: # Packet from the external network - client
19     packet_flow = LoadBalancedFlowc(h2.saddr, h2.daddr, h3.src_port, h3.dst_port, h2.npid)
20     alloc_flow_and_process_packet = False;
21     if flow_map.has(packet_flow):
22         flow_id = flow_map.get(packet_flow)
23         backend_id = flow_id_to_backend_id.get(flow_id)
24         if backend_ip_map.has_idx(backend_id):
25             flow_map.refresh_idx(flow_map.get(packet_flow), now)
26             backend = backends.get(backend_id)
27             return ([backend.nic],
28                     [ether(h1, saddr=..., daddr=backend.mac),
29                      ipv4(h2, cksum=..., daddr=backend.ip),
30                      tcpudp(h3)])
31     else:
32         flow_map.erase(packet_flow)
33         alloc_flow_and_process_packet = True
```

```
34 else:
35     alloc_flow_and_process_packet = True
36 if alloc_flow_and_process_packet:
37     if backend_ip_map.exists_with_cht(cht, _LoadBalancedFlow_hash(packet_flow)):
38         bknd = backend_ip_map.choose_with_cht(cht, _LoadBalancedFlow_hash(packet_flow))
39         if not flow_map.full():
40             idx = the_index_allocated
41             flow_map.add(packet_flow, idx, now)
42             flow_id_to_backend_id.set(idx, bknd)
43             backend = backends.get(bknd)
44             return ([backend.nic],
45                     [ether(h1, saddr=..., daddr=backend.mac),
46                      ipv4(h2, cksum=..., daddr=backend.ip),
47                      tcpudp(h3)])
48     else:
49         return ([], [])
50 else: # A heartbeat from a backend
51     bknd_addr = ip_addr(h2.saddr)
52     if backend_ip_map.has(bknd_addr):
53         backend_ip_map.refresh_idx(backend_ip_map.get(bknd_addr), now)
54     else:
55         if not backend_ip_map.full():
56             idx = the_index_allocated
57             backend_ip_map.add(bknd_addr, idx, now)
58             backends.set(idx, LoadBalancedBackendc(received_on_port, h1.saddr, h2.saddr))
59     return ([], [])
```

Push-button

Pay-as-you-go

Full-stack

Whole SW stack: framework, OS, drivers



NF: 3 KLOC

Framework, driver: 85 KLOC